

Contents

1 Structural Causal Models: A Complete Tutorial	2
1.1 What You'll Learn	2
1.2 Prerequisites	2
1.3 Table of Contents	2
1.4 1. What are Structural Causal Models?	2
1.4.1 Definition	2
1.4.2 Structural Equations	2
1.4.3 Example: Simple Linear SCM	3
1.4.4 SCMs vs Statistical Models	3
1.5 2. The Three Levels of Causation	3
1.5.1 Level 1: Association (Seeing)	3
1.5.2 Level 2: Intervention (Doing)	3
1.5.3 Level 3: Counterfactual (Imagining)	4
1.5.4 The Hierarchy	4
1.6 3. Implementing SCMs in Python	4
1.6.1 Basic SCM Class	4
1.6.2 Example: Confounded SCM	5
1.6.3 Visualizing the DAG	5
1.7 4. Interventions and the Do-Operator	5
1.7.1 Graph Surgery Interpretation	5
1.7.2 Implementing Interventions	5
1.7.3 Why Interventions Matter	6
1.7.4 Example: Confounding Bias	6
1.8 5. Counterfactual Reasoning	6
1.8.1 The Three-Step Process	6
1.8.2 Example: Coffee Counterfactual	7
1.8.3 Linear SCMs: Efficient Counterfactuals	7
1.8.4 Counterfactual vs Interventional Queries	8
1.9 6. Connection to Other Frameworks	8
1.9.1 SCMs and Potential Outcomes	8
1.9.2 SCMs and Do-Calculus	8
1.9.3 SCMs and Propensity Scores	8
1.10 7. Biological Applications	9
1.10.1 Gene Regulatory Networks	9
1.10.2 Drug Response Prediction	9
1.10.3 Cell Cycle Confounding	10
1.10.4 Perturbation Response Prediction	10
1.11 8. Advanced Topics	10
1.11.1 Identifiability of Counterfactuals	10
1.11.2 Mediation Analysis	11
1.11.3 Fairness and Discrimination	11
1.11.4 Model Explanation	11
1.12 Summary	11
1.12.1 Key Takeaways	11
1.12.2 When to Use SCMs	11
1.12.3 Further Reading	12
1.12.4 Next Steps	12
1.13 Appendix: Mathematical Details	12
1.13.1 Formal Definition of SCM	12
1.13.2 Interventions (Formal)	12
1.13.3 Counterfactuals (Formal)	12
1.13.4 Twin Network	12

1 Structural Causal Models: A Complete Tutorial

This tutorial introduces **Structural Causal Models (SCMs)**, Pearl's formal framework for causal reasoning. SCMs provide a unified approach to the three levels of causation: association, intervention, and counterfactuals.

1.1 What You'll Learn

- What structural causal models are and why they matter
- The three levels of the causal hierarchy
- How to implement SCMs in Python
- Interventions and the do-operator
- Counterfactual reasoning via abduction-action-prediction
- Applications to computational biology

1.2 Prerequisites

- Basic probability theory
- Familiarity with DAGs (see `do-calculus.md`)
- Understanding of potential outcomes (see `estimating-treatment-effects.md`)

1.3 Table of Contents

1. What are Structural Causal Models?
 2. The Three Levels of Causation
 3. Implementing SCMs in Python
 4. Interventions and the Do-Operator
 5. Counterfactual Reasoning
 6. Connection to Other Frameworks
 7. Biological Applications
 8. Advanced Topics
-

1.4 1. What are Structural Causal Models?

1.4.1 Definition

A **Structural Causal Model (SCM)** is a tuple $\mathcal{M} = \langle U, V, F \rangle$ where:

- U = **Exogenous variables** (unobserved noise, external factors)
- V = **Endogenous variables** (observed variables in the system)
- F = **Structural equations** (functions defining how V are generated from U and other V)

1.4.2 Structural Equations

Each endogenous variable V_i is determined by a structural equation:

$$V_i := f_i(\text{PA}_i, U_i)$$

where: * PA_i are the parents of V_i (other endogenous variables) * U_i is the exogenous noise for V_i * f_i is a deterministic function

Key insight: The $:=$ symbol means “is determined by” (not “equals”). This is an **assignment**, not an algebraic equation.

1.4.3 Example: Simple Linear SCM

Consider a simple causal relationship: smoking (X) causes lung cancer (Y).

Structural equations:

$$X := U_X$$

$$Y := 2X + U_Y$$

where: * $U_X \sim \mathcal{N}(0, 1)$ (individual propensity to smoke) * $U_Y \sim \mathcal{N}(0, 0.5)$ (other factors affecting cancer)

Interpretation: * Smoking is determined by individual propensity * Cancer risk is determined by smoking (coefficient 2) plus other factors

1.4.4 SCMs vs Statistical Models

Aspect	Statistical Model	Structural Causal Model
Focus	Associations, predictions	Causal mechanisms
Equations	$Y = 2X + \epsilon$	$Y := 2X + U_Y$
Interpretation	Correlation	Causation
Interventions	Not defined	Well-defined (do-operator)
Counterfactuals	Not computable	Computable

The key difference: SCMs model the **data-generating process**, not just the data distribution.

1.5 2. The Three Levels of Causation

Pearl's **Ladder of Causation** describes three increasingly powerful types of causal reasoning:

1.5.1 Level 1: Association (Seeing)

Question: What is?

Query: $P(Y | X)$

Example: "What is the cancer rate among smokers?"

Computation: Observational data is sufficient

$$P(Y | X) = \frac{P(X, Y)}{P(X)}$$

Limitation: Cannot distinguish causation from confounding

1.5.2 Level 2: Intervention (Doing)

Question: What if we do?

Query: $P(Y | do(X))$

Example: "What would the cancer rate be if we forced everyone to smoke?"

Computation: Requires causal assumptions (DAG + do-calculus)

$$P(Y | do(X)) \neq P(Y | X) \text{ (in general)}$$

Key insight: Interventions break incoming causal arrows

1.5.3 Level 3: Counterfactual (Imagining)

Question: What if we had done?

Query: $P(Y_x | X', Y')$

Example: “Would this patient have survived if they had not smoked, given that they smoked and died?”

Computation: Requires full SCM (structural equations)

Key insight: Counterfactuals are **individual-level** statements, not population-level

1.5.4 The Hierarchy

Level 3: Counterfactuals

```
↑
| (requires SCM)
|
```

Level 2: Interventions

```
↑
| (requires DAG)
|
```

Level 1: Associations

```
↑
| (requires data)
```

Important: You cannot answer Level 3 questions with only Level 2 tools, and you cannot answer Level 2 questions with only Level 1 tools.

1.6 3. Implementing SCMs in Python

1.6.1 Basic SCM Class

```
from causalbiolab.scm import StructuralCausalModel, SCMVariable
from scipy import stats

# Define variables
variables = {
    'X': SCMVariable(
        name='X',
        equation=lambda u_x: u_x,
        parents=[],
        noise_dist=stats.norm(0, 1)
    ),
    'Y': SCMVariable(
        name='Y',
        equation=lambda x, u_y: 2*x + u_y,
        parents=['X'],
        noise_dist=stats.norm(0, 0.5)
    )
}
```

```

# Create SCM
scm = StructuralCausalModel(variables)

# Sample observational data
data = scm.sample(n_samples=1000)

```

1.6.2 Example: Confounded SCM

```

#  $Z \rightarrow X, Z \rightarrow Y$  ( $Z$  is confounder)
variables = {
    'Z': SCMVariable(
        name='Z',
        equation=lambda u_z: u_z,
        parents=[],
        noise_dist=stats.norm(0, 1)
    ),
    'X': SCMVariable(
        name='X',
        equation=lambda z, u_x: z + u_x,
        parents=['Z'],
        noise_dist=stats.norm(0, 0.5)
    ),
    'Y': SCMVariable(
        name='Y',
        equation=lambda x, z, u_y: 2*x + z + u_y,
        parents=['X', 'Z'],
        noise_dist=stats.norm(0, 0.5)
    )
}

```

```

scm_conounded = StructuralCausalModel(variables)

```

1.6.3 Visualizing the DAG

Every SCM induces a DAG:

```

dag = scm_conounded.get_dag()
# {'Z': ['X', 'Y'], 'X': ['Y'], 'Y': []}

```

1.7 4. Interventions and the Do-Operator

1.7.1 Graph Surgery Interpretation

The do-operator $do(X = x)$ means:

1. Cut all incoming edges to X in the DAG
2. Set $X = x$ (constant value)

This creates a **mutilated graph** $G_{\bar{X}}$.

1.7.2 Implementing Interventions

```

# Intervene: set X = 1.5
scm_do_x = scm.intervene({'X': 1.5})

```

```

# Sample from intervened distribution
data_do_x = scm_do_x.sample(1000)

# Compare observational vs interventional
print(f"E[Y | X=1.5] = {data['Y'][data['X'] > 1.4].mean():.3f}") # Observational
print(f"E[Y | do(X=1.5)] = {data_do_x['Y'].mean():.3f}") # Interventional

```

1.7.3 Why Interventions Matter

Observational:

$$P(Y | X) = \sum_Z P(Y | X, Z)P(Z | X)$$

Includes confounding through Z .

Interventional:

$$P(Y | do(X)) = \sum_Z P(Y | X, Z)P(Z)$$

Removes confounding (no $P(Z | X)$).

1.7.4 Example: Confounding Bias

```

# Observational: biased by Z
data_obs = scm_confounded.sample(1000)
obs_effect = data_obs['Y'][data_obs['X'] > 0].mean() - data_obs['Y'][data_obs['X'] <= 0].mean()

# Interventional: unbiased
scm_do_x1 = scm_confounded.intervene({'X': 1.0})
scm_do_x0 = scm_confounded.intervene({'X': 0.0})
data_do_x1 = scm_do_x1.sample(1000)
data_do_x0 = scm_do_x0.sample(1000)
causal_effect = data_do_x1['Y'].mean() - data_do_x0['Y'].mean()

print(f"Observational effect: {obs_effect:.3f}") # Biased
print(f"Causal effect: {causal_effect:.3f}") # True effect 2.0

```

1.8 5. Counterfactual Reasoning

1.8.1 The Three-Step Process

Counterfactual computation follows Pearl's **abduction-action-prediction** framework:

1.8.1.1 Step 1: Abduction Infer the exogenous variables U from observed data.

Given observed $(X = x, Y = y)$, solve for U :

$$U_X = x$$

$$U_Y = y - 2x$$

1.8.1.2 Step 2: Action Modify the SCM by applying intervention $do(X = x')$.

Create mutilated SCM with $X := x'$ (constant).

1.8.1.3 Step 3: Prediction Compute the query variable using the inferred U and modified SCM.

$$Y_{x'} = 2x' + U_Y = 2x' + (y - 2x)$$

1.8.2 Example: Coffee Counterfactual

Scenario: Joe drank coffee ($T = 1$) and stayed awake ($Y = 1$). Would he have stayed awake if he hadn't drunk coffee?

SCM:

$$Y := T \cdot U + (T - 1)(U - 1)$$

Computation:

```
from causalbiolab.scm.counterfactuals import compute_counterfactual

# Observed: T=1, Y=1
# Query: What if T=0?
y_cf = compute_counterfactual(
    scm,
    observed={'T': 1, 'Y': 1},
    intervention={'T': 0},
    query='Y'
)

print(f"Counterfactual outcome: {y_cf} # Y_0 = 0")
```

Interpretation: Joe would not have stayed awake without coffee.

1.8.3 Linear SCMs: Efficient Counterfactuals

For linear SCMs, counterfactuals are particularly simple:

```
from causalbiolab.scm.counterfactuals import LinearSCM

# Linear SCM: X -> Y with Y = 2X + U_Y
scm_linear = LinearSCM(
    coefficients={'Y': {'X': 2.0}},
    noise_distributions={'X': stats.norm(0, 1), 'Y': stats.norm(0, 0.5)}
)

# Counterfactual: observed X=1, Y=3; what if X=2?
y_cf = scm_linear.counterfactual(
    observed={'X': 1, 'Y': 3},
    intervention={'X': 2},
    query='Y'
)

print(f"Y_{X=2} = {y_cf:.3f} # Should be 5.0")
```

1.8.4 Counterfactual vs Interventional Queries

Query Type	Question	Requires
Interventional	$E[Y \mid do(X = x)]$	Population-level, DAG sufficient
Counterfactual	Y_x for individual	Individual-level, full SCM needed

Key difference: Counterfactuals use **observed** values to infer individual-specific U , then predict under intervention.

1.9 6. Connection to Other Frameworks

1.9.1 SCMs and Potential Outcomes

Potential outcomes framework (Rubin): * $Y_i(1), Y_i(0)$ are potential outcomes * ATE = $E[Y(1) - Y(0)]$

SCM perspective: * Potential outcomes are **counterfactuals** * $Y_i(1) = f_Y(\text{PA}_i, U_i)$ when $T_i := 1$ * $Y_i(0) = f_Y(\text{PA}_i, U_i)$ when $T_i := 0$

Connection:

$$\text{ATE} = E[Y \mid do(T = 1)] - E[Y \mid do(T = 0)]$$

SCMs provide the **mechanism** underlying potential outcomes.

1.9.2 SCMs and Do-Calculus

Do-calculus provides rules for identifying $P(Y \mid do(X))$ from observational data.

SCMs provide the **implementation:** * DAG structure comes from SCM * Do-operator is graph surgery on SCM * Identification formulas compute expectations in mutilated SCM

Example: Back-door adjustment

$$P(Y \mid do(X)) = \sum_Z P(Y \mid X, Z)P(Z)$$

In SCM terms: 1. Sample Z from marginal (unaffected by intervention) 2. Sample Y from conditional given X and Z

1.9.3 SCMs and Propensity Scores

Propensity score: $e(X) = P(T = 1 \mid X)$

In SCM: * $e(X)$ emerges from structural equation for T * IPW reweights to simulate $do(T = t)$ * SCM makes explicit what IPW assumes

1.10 7. Biological Applications

1.10.1 Gene Regulatory Networks

SCM:

$$\text{TF} := U_{\text{TF}}$$

$$\text{Gene} := \sigma(\text{TF}) + U_{\text{Gene}}$$

$$\text{Protein} := \text{Gene} \cdot \exp(U_{\text{Protein}})$$

where $\sigma(x) = 1/(1 + e^{-x})$ is sigmoid activation.

Questions: * **Intervention:** What if we knock out the TF? $do(\text{TF} = 0)$ * **Counterfactual:** Would this cell express the gene if TF was higher?

```
from causalbiolab.scm.examples import gene_regulation_scm

scm_gene = gene_regulation_scm()

# Intervention: knockout TF
scm_knockout = scm_gene.intervene({'TF': 0})
data_knockout = scm_knockout.sample(1000)

print(f"Gene expression under knockout: {data_knockout['Gene'].mean():.3f}")
```

1.10.2 Drug Response Prediction

SCM:

$$\text{Genotype} := U_G > 0$$

$$\text{DrugMetabolism} := 0.5 \cdot \text{Genotype} \cdot \text{Dose} + U_M$$

$$\text{Response} := 2 \cdot \text{Dose} - \text{DrugMetabolism} + U_R$$

Counterfactual question: “Would this patient respond better with a different genotype?”

```
from causalbiolab.scm.examples import drug_response_scm

scm_drug = drug_response_scm()

# Observed: poor metabolizer (Genotype=1), low response
observed = {'Genotype': 1, 'DrugDose': 1.0, 'Response': 1.5}

# Counterfactual: what if normal metabolizer?
response_cf = scm_drug.counterfactual(
    observed=observed,
    intervention={'Genotype': 0},
    query='Response'
)
```

```
print(f"Counterfactual response: {response_cf:.3f}")
```

1.10.3 Cell Cycle Confounding

SCM:

$$\text{CellCycle} := U_{CC}$$

$$\text{Transfection} := \sigma(\text{CellCycle}) + U_T$$

$$\text{GeneExpression} := 2 \cdot \text{Transfection} + 0.5 \cdot \text{CellCycle} + U_G$$

Intervention: What if we control for cell cycle?

```
from causalbiolab.scm.examples import cell_cycle_confounding_scm

scm_cc = cell_cycle_confounding_scm()

# Observational: confounded
data_obs = scm_cc.sample(1000)

# Interventional: fix cell cycle
scm_fixed_cc = scm_cc.intervene({'CellCycle': 0})
data_fixed = scm_fixed_cc.sample(1000)

# Compare transfection effect
print("Observational correlation:", np.corrcoef(data_obs['Transfection'], data_obs['GeneExpression']))[0]
print("Causal effect (CC fixed):", np.corrcoef(data_fixed['Transfection'], data_fixed['GeneExpression']))[0]
```

1.10.4 Perturbation Response Prediction

Use case: Predict phenotype after CRISPR knockout

SCM approach: 1. Learn SCM from observational single-cell data 2. Intervene on target gene: $do(\text{Gene} = 0)$ 3. Predict downstream effects

Advantage over black-box models: Mechanistic interpretation, compositionality for multi-gene perturbations

1.11 8. Advanced Topics

1.11.1 Identifiability of Counterfactuals

Question: Can we compute counterfactuals from data?

Answer: Depends on the SCM structure.

Identifiable cases: * Linear SCMs with Gaussian noise * Monotonic functions with specific noise distributions * Discrete variables with finite support

Non-identifiable cases: * Nonlinear SCMs with arbitrary noise * Hidden confounders between treatment and outcome

Practical implication: For biology, often need to make **parametric assumptions** about structural equations.

1.11.2 Mediation Analysis

Question: How much of the effect goes through mediator M ?

Natural Direct Effect (NDE):

$$\text{NDE} = E[Y_{X=1, M=M_0} - Y_{X=0, M=M_0}]$$

Natural Indirect Effect (NIE):

$$\text{NIE} = E[Y_{X=1, M=M_1} - Y_{X=1, M=M_0}]$$

where M_t is the mediator value under $X = t$.

Requires: Counterfactual reasoning (Level 3)

1.11.3 Fairness and Discrimination

Counterfactual fairness: A decision is fair if:

$$P(\hat{Y}_A \mid X, A = a) = P(\hat{Y}_{A'} \mid X, A = a)$$

for all a, a' (protected attributes).

Interpretation: Outcome would be the same if individual had different protected attribute.

Application: Ensure drug recommendations don't discriminate based on race/gender.

1.11.4 Model Explanation

Counterfactual explanations: "Your loan was denied because if your income were \$10K higher, it would have been approved."

SCM approach: 1. Learn SCM from data 2. Compute counterfactuals for feature changes 3. Find minimal changes that flip prediction

1.12 Summary

1.12.1 Key Takeaways

1. SCMs formalize causation through structural equations
2. Three levels of causation require increasingly strong assumptions
3. Interventions break causal arrows (do-operator)
4. Counterfactuals require abduction-action-prediction
5. SCMs unify potential outcomes, do-calculus, and graphical models

1.12.2 When to Use SCMs

Use SCMs when you need: * Individual-level predictions (counterfactuals) * Mechanistic understanding (not just associations) * Composition of interventions (multi-gene knockouts) * Explanation of model predictions

Don't use SCMs when: * Only population-level effects needed (use potential outcomes) * Only identification needed (use do-calculus) * Structural equations unknown (use nonparametric methods)

1.12.3 Further Reading

- Pearl, J. (2009). *Causality: Models, Reasoning, and Inference*
- Pearl, J., & Mackenzie, D. (2018). *The Book of Why*
- Peters, J., Janzing, D., & Schölkopf, B. (2017). *Elements of Causal Inference*

1.12.4 Next Steps

1. **Interactive notebook:** Work through examples hands-on
 2. **Biological applications:** Apply to gene networks, drug response
 3. **Integration:** Connect SCMs to existing causal inference tools
-

1.13 Appendix: Mathematical Details

1.13.1 Formal Definition of SCM

An SCM $\mathcal{M} = \langle U, V, F, P(U) \rangle$ consists of:

- $U = \{U_1, \dots, U_m\}$: exogenous variables
- $V = \{V_1, \dots, V_n\}$: endogenous variables
- $F = \{f_1, \dots, f_n\}$: structural functions where $V_i = f_i(\text{PA}_i, U_i)$
- $P(U)$: joint distribution over exogenous variables

1.13.2 Interventions (Formal)

The **mutilated model** $\mathcal{M}_{\bar{X}}$ under $do(X = x)$ is:

$$\mathcal{M}_{\bar{X}} = \langle U, V, F_{\bar{X}}, P(U) \rangle$$

where $F_{\bar{X}}$ replaces f_X with constant function $f_X(\cdot) = x$.

1.13.3 Counterfactuals (Formal)

The **counterfactual** $Y_x(u)$ is the value of Y in model $\mathcal{M}_{\bar{X}}$ with exogenous values $U = u$:

$$Y_x(u) = f_Y^{\mathcal{M}_{\bar{X}}}(\text{PA}_Y, U_Y)$$

evaluated recursively in topological order.

1.13.4 Twin Network

Counterfactuals can be visualized as a **twin network**: * Factual world: actual observations * Counterfactual world: intervened model * Shared exogenous variables U link the two worlds

This explains why counterfactuals are individual-specific: they depend on the specific U for that individual.